

Summary

These authors examine the relationship between take-up of transfers and lifetime income and consumption using PSID data. They find that take-up is negatively correlated with lifetime income and consumption across households, holding annual income fixed. They use a theoretical model that maps these empirical findings to the value of making take-up of transfers automatic, which is akin to mandating take-up among all eligible households.

This paper presents a thought-provoking and well-executed set of empirical results and some interesting theoretical reasoning. However, in my view, some implicit structure in the model leads the authors to over-state the policy implications of their results.

Comment

Let me summarize the part of the argument with which I agree. For the sake of illustration, suppose conditional on annual incomes, there are two types of people who are eligible for a transfer. One type just suffered a negative shock and the other is long-term eligible. The first type is the typical target for *social insurance*, and the other is the typical target for *redistribution*. The main empirical result of the paper is that making transfers automatic primarily benefits the insurance type, and in general the insurance type has higher lifetime income and consumption than the redistribution type. So, the redistributive value of a transfer to the redistribution type is higher than that of the insurance type.

How we map this to the conclusion about non-optimality of automatic payments requires a few more premises with which I do not agree. Making these assumptions more transparent and relaxing them where appropriate would allow for a more nuanced and credible discussion of the policy implications of the empirical findings. Here are the implicit assumptions with which I am struggling:

- A1) it is infeasible to levy a tax on lifetime income or consumption,
- A2) transfer programs do not insure uninsurable risk,
- A3) the transfer program's budget must be held fixed, so reforms to automatic transfers must be financed by changes to benefit levels.

A1 motivates the structure used to weigh the value of a dollar to redistributive types versus insurance types. If to the contrary lifetime income or consumption taxation were feasible, the Atkinson-Stiglitz theorem (or the Kaplow 2011 argument the authors cite) suggests that at the global optimum, the lifetime income tax handles redistribution and other transfers are mainly justified by their social insurance value. In this case, we are likely to find, based on the empirical data, that at the global optimum, automating transfers is welfare-improving, because it helps target the insurance types. I do not have a problem with A1) on its own, but it should be more transparent (e.g. in the introduction) that this is how you are thinking about welfare.

I do disagree with A2). A textbook rationale for many transfer programs is that individuals have difficulty insuring major risks due to market failures. Assuming that individuals can perfectly smooth consumption, so that welfare weights should be based on lifetime income/consumption, assumes this rationale away. This in effect focuses the paper's evaluation of transfers on redistribution rather than the *efficiency* gains from insuring risk. Looking at life-time consumption or income is a standard way to value redistribution, but that does not mean that the efficiency benefits of providing social insurance do not exist. One way to quantify the value of these insurance benefits is to examine the evolution of consumption around shocks, within households (Landais and Spinnewijn 2021 *REStud*). I note that the findings in this paper and other related papers suggest that insurance values are substantial. As a minor comment, I note that this paper is also part of a literature leveraging measures of consumption in administrative data to evaluate redistributive and insurance aspects of tax and transfer policies. Doing so is feasible in many European countries where the limitations of administrative data the authors describe are not present. Some citations are in order.

Finally, A3) justifies the budget neutrality of the reform the paper considers. With a broader notion of budget neutrality that reflects the fungibility of money, the question we need to answer to evaluate a reform to automatic take-up would be whether the (welfare-weighted) value of transfers to the insurance types exceeds the fiscal cost of the reform, where the latter is weighed according to the next-best use of government funds (the MVPF). When we focus on redistribution, the empirical results suggest the marginal recipient values transfers less than the average recipient, but whether the benefit to the marginal recipient exceeds the cost of providing transfers to the marginal recipient is unclear. Forcing a tradeoff between the average and marginal recipient via A3 yields an unambiguous policy implication, but this tradeoff does not exist if public funds are fungible. If we additionally relax A2), the value of transfers to the insurance types goes up, which also works in favor of automation. In the end, I'm not sure we have enough data to make a compelling case about whether automation is welfare-improving, but I do think the paper tells an interesting story about selection in to transfers.